

Interpreting logistic models

Frank Edwards

- Remember that logistic regression is a GLM with a logit link function

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$
- For logistic regression, g is the logit function: $\text{logit}(y) = \mathbf{X}\beta$

- Remember that logistic regression is a GLM with a logit link function
- A GLM takes the form: $g(y) = \mathbf{X}\beta$
- For logistic regression, g is the logit function: $\text{logit}(y) = \mathbf{X}\beta$

Let's return to the grad school admission example

```
admits <- read_csv("./data/binary.csv")  
summary(admits)
```

##	admit	gre	gpa	rank
##	Min. :0.0000	Min. :220.0	Min. :2.260	Min. :1.000
##	1st Qu.:0.0000	1st Qu.:520.0	1st Qu.:3.130	1st Qu.:2.000
##	Median :0.0000	Median :580.0	Median :3.395	Median :2.000
##	Mean :0.3175	Mean :587.7	Mean :3.390	Mean :2.485
##	3rd Qu.:1.0000	3rd Qu.:660.0	3rd Qu.:3.670	3rd Qu.:3.000
##	Max. :1.0000	Max. :800.0	Max. :4.000	Max. :4.000

Let's look at this as the distribution of the probability of admissions across the data

- First, fit an intercept-only logistic regression model

```
m0 <- stan_glm(admit ~ 1, data = admits, family = "binomial", refresh = 0)
m0
```

```
## stan_glm
## family:      binomial [logit]
## formula:     admit ~ 1
## observations: 400
## predictors:  1
## -----
##              Median MAD_SD
## (Intercept) -0.8    0.1
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

- What does this model tell us?

What does this model tell us?

```
coef(m0) ## log odds
```

```
## (Intercept)  
## -0.7697838
```

```
exp(coef(m0)) ## odds
```

```
## (Intercept)  
## 0.4631132
```

```
exp(coef(m0))/(1 + exp(coef(m0))) ## probability
```

```
## (Intercept)  
## 0.3165259
```

```
mean(admits$admit) ## mean proportion admitted
```

```
## [1] 0.3175
```


Let's add a predictor

```
m1 <- stan_glm(admit ~ gre, data = admits, family = "binomial", refresh = 0)
```


```
m1
```

```
## stan_glm
## family:      binomial [logit]
## formula:      admit ~ gre
## observations: 400
## predictors:    2
## -----
##              Median MAD_SD
## (Intercept) -2.9      0.6
## gre          0.0      0.0
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

Hmmmm. why is β_1 so small?

Posterior distribution of beta_1: What does this tell us about GRE scores and admission?

```
mcmc_dens(m1, pars = "gre")
```



To ease interpretation, let's scale GRE

Scale mean-centers and SD scales variables: $\text{scale}(x_i) = \frac{x_i - \bar{x}}{sd(x)}$

```
m2 <- stan_glm(admit ~ scale(gre), data = admits, family = "binomial", refresh = 0)
```

Interpret the model

```
coef(m2)
```

```
## (Intercept)  scale(gre)  
##   -0.7968293    0.4154692
```

```
sd(admits$gre)
```

```
## [1] 115.5165
```

Remember: $\text{logit}(p) = \log\left(\frac{p}{1-p}\right) = \mathbf{X}\beta$

And: $\text{logit}^{-1}(\mathbf{X}\beta) = \frac{\exp(\mathbf{X}\beta)}{\exp(\mathbf{X}\beta)+1}$

Interpret the model

```
coef(m2)
```

```
## (Intercept)  scale(gre)  
##   -0.7968293    0.4154692
```

```
sd(admits$gre)
```

```
## [1] 115.5165
```

Remember: $\text{logit}(p) = \log\left(\frac{p}{1-p}\right) = \mathbf{X}\beta$

And: $\text{logit}^{-1}(\mathbf{X}\beta) = \frac{\exp(\mathbf{X}\beta)}{\exp(\mathbf{X}\beta)+1}$

- What does β_0 mean?

Interpret the model

```
coef(m2)
```

```
## (Intercept)  scale(gre)  
##   -0.7968293    0.4154692
```

```
sd(admits$gre)
```

```
## [1] 115.5165
```

Remember: $\text{logit}(p) = \log\left(\frac{p}{1-p}\right) = \mathbf{X}\beta$

And: $\text{logit}^{-1}(\mathbf{X}\beta) = \frac{\exp(\mathbf{X}\beta)}{\exp(\mathbf{X}\beta)+1}$

- What does β_0 mean?
- What does β_1 mean?

$$e^{y_1+y_2} = e^{y_1} e^{y_2}$$

$$e^{y_1+y_2} = e^{y_1} e^{y_2}$$

and

$$e^{y_1-y_2} = \frac{e^{y_1}}{e^{y_2}}$$

$$e^{y_1+y_2} = e^{y_1} e^{y_2}$$

and

$$e^{y_1-y_2} = \frac{e^{y_1}}{e^{y_2}}$$

We can take our logistic regression:

$$\log \left(\frac{p}{1-p} \right) = \beta_0 + \beta_1 x_1$$

We can take our logistic regression:

$$\log \left(\frac{p}{1-p} \right) = \beta_0 + \beta_1 x_1$$

Exponentiate both sides and we obtain

$$\left(\frac{p}{1-p} \right) = e^{\beta_0 + \beta_1 x_1}$$

A funny thing happens with our linear predictor when exponentiated

$$\left(\frac{p}{1-p} \right) = e^{\beta_0 + \beta_1 x_1} = e^{\beta_0} e^{\beta_1 x_1}$$

A funny thing happens with our linear predictor when exponentiated

$$\left(\frac{p}{1-p} \right) = e^{\beta_0 + \beta_1 x_1} = e^{\beta_0} e^{\beta_1 x_1}$$

The odds of p are related to our predictors through a multiplicative, rather than additive relationship

Odds are defined as the probability of the event occurring divided by the probability of probability of the event not occurring. To obtain odds in a logistic regression, we exponentiate both sides:

$$\frac{p}{1-p} = e^{\beta_0 + \beta_1 x_1}$$

Odds are defined as the probability of the event occurring divided by the probability of probability of the event not occurring. To obtain odds in a logistic regression, we exponentiate both sides:

$$\frac{p}{1-p} = e^{\beta_0 + \beta_1 x_1}$$

The odds ratio is the ratio of two odds - or the proportional change in odds. We can obtain an isolated estimate for the relationship between $\beta_1 x_{1i}$ and y this way:

$$\frac{\text{Odds}(p|x_1 = 1)}{\text{Odds}(p|x_1 = 0)} = \frac{e^{\beta_0 + \beta_1}}{e^{\beta_0}} = \frac{e^{\beta_0} \times e^{\beta_1}}{e^{\beta_0}} = e^{\beta_1}$$

The odds ratio can be interpreted as the change in odds of p for a one-unit change in x_1 .

- Odds ratios appear convenient - e^{β_1} is a percent change in $p/(1 - p)$ for a one-unit change in x_1

- Odds ratios appear convenient - e^{β_1} is a percent change in $p/(1 - p)$ for a one-unit change in x_1

How do they work?

In our example

```
coef(m2)
```

```
## (Intercept)  scale(gre)  
##   -0.7968293    0.4154692
```

```
exp(coef(m2))
```

```
## (Intercept)  scale(gre)  
##    0.4507559    1.5150815
```

What does β_1 mean? What does e^{β_1} mean?

The odds of admission are 1.52 times higher for a student with a GRE score one standard deviation above the mean than they are for a student with a mean GRE score.

The odds of admission are 1.52 times higher for a student with a GRE score one standard deviation above the mean than they are for a student with a mean GRE score.

Any trouble you can anticipate here?

Expected changes for a 1 SD increase in GRE scores

```
mean(admits$gre)
```

```
## [1] 587.7
```

```
sd(admits$gre)
```

```
## [1] 115.5165
```

```
# log odds increase for +1 SD GRE
```

```
coef(m2)[2]
```

```
## scale(gre)
```

```
## 0.4154692
```

```
# odds proportional change for +1 SD GRE
```

```
exp(coef(m2)[2])
```

```
## scale(gre)
```

```
## 1.515082
```

On the probability scale

```
# change in probability for a low GRE + 1
p_low <- predict(m2, newdata = data.frame(gre = 200), type = "response")
p_lowP1 <- predict(m2, newdata = data.frame(gre = 200 + sd(admits$gre)), type = "response")

p_lowP1 - p_low

##           1
## 0.04228963
```

On the probability scale

```
# change in probability for a mid GRE + 1
p_mid <- predict(m2, newdata = data.frame(gre = 500), type = "response")
p_midP1 <- predict(m2, newdata = data.frame(gre = 500 + sd(admits$gre)), type = "response")

p_midP1 - p_mid

##          1
## 0.08470265
```


On the probability scale

```
# change in probability for a mid GRE + 1
p_hi <- predict(m2, newdata = data.frame(gre = 650), type = "response")
p_hiP1 <- predict(m2, newdata = data.frame(gre = 650 + sd(admits$gre)), type = "response")


p_hiP1 - p_hi

##          1
## 0.1003912
```

Bayesian inference with logistic regression

Posterior distributions of Beta parameters

```
library(bayesplot)  
mcmc_hist(m2)
```



To interpret, let's use the linear predictor

Rather than interpret β_1 , let's look at the linear relationship directly

```
library(tidybayes)
m2_lp <- linpred_draws(m2, newdata = admits, ndraws = 1000)


head(m2_lp)
```



```
## # A tibble: 6 x 9
## # Groups:   admit, gre, gpa, rank, .row [1]
##   admit gre gpa rank .row .chain .iteration .draw .linpred
##   <dbl> <dbl> <dbl> <dbl> <int> <int> <int> <int> <dbl>
## 1     0  380  3.61    3     1    NA      NA     1  -1.09
## 2     0  380  3.61    3     1    NA      NA     2  -1.24
## 3     0  380  3.61    3     1    NA      NA     3  -1.69
## 4     0  380  3.61    3     1    NA      NA     4  -1.56
## 5     0  380  3.61    3     1    NA      NA     5  -1.17
## 6     0  380  3.61    3     1    NA      NA     6  -1.51
```


To interpret, let's use the linear predictor

```
ggplot(m2_lp, aes(x = gre, y = .linpred)) + stat_lineribbon() + scale_fill_brewer()
```



How about on the odds scale

```
ggplot(m2_lp, aes(x = gre, y = exp(.linpred))) + stat_lineribbon() + scale_fill_brewer()
```



We can also estimate expected probability

```
m2_ep <- epred_draws(m2, newdata = admits, ndraws = 1000)
```

```
head(m2_ep)
```

```
## # A tibble: 6 x 9
```

```
## # Groups:   admit, gre, gpa, rank, .row [1]
```

##	admit	gre	gpa	rank	.row	.chain	.iteration	.draw	.epred
##	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<dbl>
## 1	0	380	3.61	3	1	NA	NA	1	0.201
## 2	0	380	3.61	3	1	NA	NA	2	0.185
## 3	0	380	3.61	3	1	NA	NA	3	0.172
## 4	0	380	3.61	3	1	NA	NA	4	0.180
## 5	0	380	3.61	3	1	NA	NA	5	0.221
## 6	0	380	3.61	3	1	NA	NA	6	0.199

And visualized

```
ggplot(m2_ep, aes(x = gre, y = .epred)) + stat_lineribbon() + scale_fill_brewer()
```

