

# Count data and the Poisson distribution

---

Frank Edwards

- Counts are cumulative totals of the number of incidences of some event, generally across time or place

- Counts are cumulative totals of the number of incidences of some event, generally across time or place
- Counts are positive integers  $\in [0, \infty]$

- Counts are cumulative totals of the number of incidences of some event, generally across time or place
- Counts are positive integers  $\in [0, \infty]$

## Counts as extensions of binary data

- Counts can be thought of as repeated binary trials
- $\sum y_i$  where  $y$  is equal to 1 or 0 provides a count
- Generally, we could treat `sum(y==1) + sum(y==0)` or `nrow(y)` as the exposure, or denominator for a rate. Why?

# Data for today: National Women's Soccer League stats

```
## data from https://github.com/adr1n1/nwslR
# devtools::install_github("adr1n1/nwslR")
library(nwslR)
data("player")
data("fieldplayer_overall_season_stats")
head(player, n=2)
```

```
## # A tibble: 2 x 5
##   person_id player      nation pos  name_other
##   <dbl> <chr>      <chr> <chr> <chr>
## 1     342 Marisa Abegg  USA   DF   <NA>
## 2     117 Danesha Adams USA   FW,MF <NA>
```

```
head(fieldplayer_overall_season_stats, n=2)
```

```
## # A tibble: 2 x 14
##   person_id season nation pos  team_id  mp starts  min  gls  ast  pk
##   <int> <dbl> <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     342  2013 USA   DF   WAS      5     4   NA     0     0     0
## 2     117  2013 USA   FW,MF NJ     20    20   NA     3     3     1
## # i 3 more variables: p_katt <dbl>, crd_y <dbl>, crd_r <dbl>
```

```
# check the help files with ?(fieldplayer_overall_season_stats) for codebook
```

# make a joined table with players names

```
### attaching names
```

```
dat<-fieldplayer_overall_season_stats %>%  
  left_join(player)
```

```
glimpse(dat)
```

```
## Rows: 1,350  
## Columns: 16  
## $ person_id <dbl> 342, 117, 6, 300, 202, 202, 28, 290, 56, 313, 363, 454, 414~  
## $ season <dbl> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013,~  
## $ nation <chr> "USA", "USA", "ESP", "USA", "USA", "USA", "USA", "USA", "US~  
## $ pos <chr> "DF", "FW,MF", "FW", "DF,MF", "DF", "DF", "DF", "DF", "MF",~  
## $ team_id <chr> "WAS", "NJ", "WNY", "KC", "POR", "BOS", "WNY", "SEA", "NJ",~  
## $ mp <dbl> 5, 20, 15, 22, 4, 11, 7, 22, 7, 20, 3, 2, 1, 22, 20, 6, 1, ~  
## $ starts <dbl> 4, 20, 14, 22, 2, 11, 2, 22, 5, 11, 0, 1, 0, 22, 16, 5, 0, ~  
## $ min <dbl> NA, NA, 3, 1900, 212, 990, NA, NA, NA, NA, NA, 123, NA, 198~  
## $ gls <dbl> 0, 3, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 3, 2, 0, 0, 5, 0,~  
## $ ast <dbl> 0, 3, NA, 5, 0, 1, NA, 1, 0, 2, 0, 0, 0, 2, 1, 0, 0, 0, 4, ~  
## $ pk <dbl> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~  
## $ p_katt <dbl> 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~  
## $ crd_y <dbl> 0, 0, 2, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 3, 0,~  
## $ crd_r <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~  
## $ player <chr> "Marisa Abegg", "Danesha Adams", "Adriana", "Leigh Ann Brow~  
## $ name_other <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
```

## Approaches to modeling count data

---



# The Poisson model

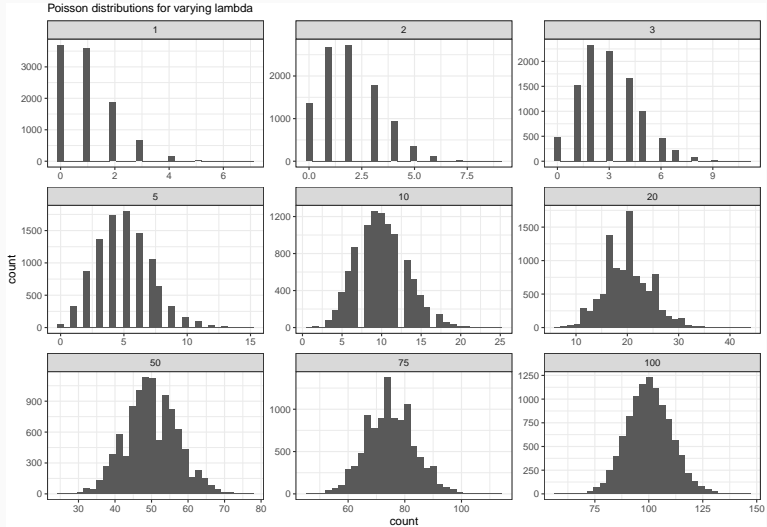
Where  $y$  is a non-negative integer (count)

$$y \sim \text{Poisson}(\lambda)$$

$$E(y) = \bar{y} = \lambda$$

$$\text{Var}(y) = \lambda$$

# The Poisson Distribution



## Let's look at each Poisson variable

```
pois_demo %>% group_by(lambda) %>%  
  summarise(mean = mean(count),  
            variance = var(count))
```

```
## # A tibble: 9 x 3  
##   lambda    mean variance  
##   <dbl> <dbl>   <dbl>  
## 1      1    1.01     1.01  
## 2      2    2.00     1.99  
## 3      3    2.97     2.94  
## 4      5    5.01     5.00  
## 5     10   10.0    10.0  
## 6     20   20.0    20.3  
## 7     50   49.9    50.0  
## 8     75   75.0    72.8  
## 9    100  100.    101.
```

For a count variable  $y$ , we can specify a Poisson GLM with a log link function

$$y \sim \text{Poisson}(\lambda)$$

$$\lambda = e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}$$

For a count variable  $y$ , we can specify a Poisson GLM with a log link function

$$y \sim \text{Poisson}(\lambda)$$

$$\lambda = e^{\beta_0 + \beta_1 x_1 \cdots \beta_n x_n}$$

What is  $\log(\lambda)$  equal to?

$$E(y|x) = e^{\lambda}$$

$$\log(E(y|x)) = \lambda = x\beta$$

$$E(y|x) = e^{\lambda}$$

$$\log(E(y|x)) = \lambda = x\beta$$

if a GLM is defined as  $g(\mu) = x\beta$  with link function  $g$ , what is the link function for the Poisson GLM?

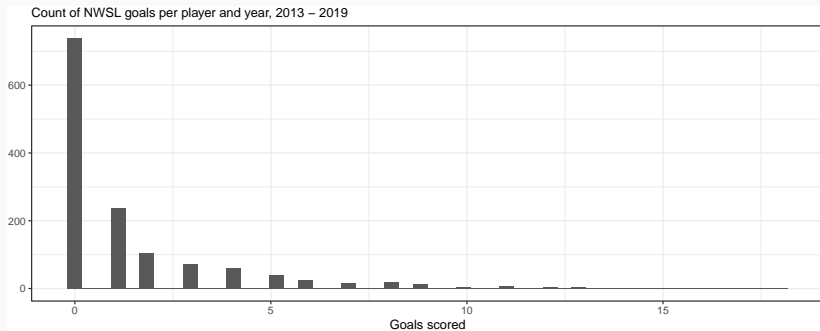
## Modeling NWSL data using a Poisson GLM

---



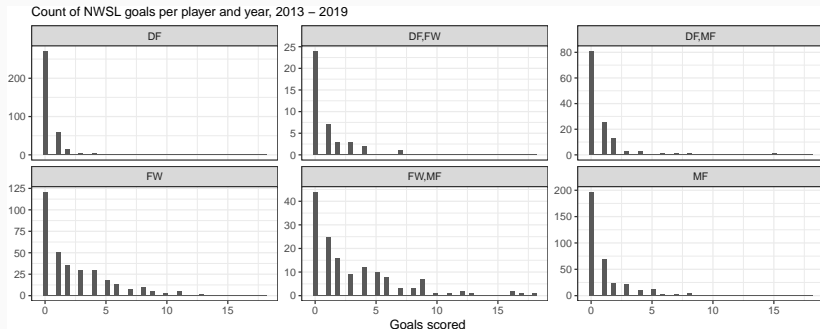
# Goal scoring

```
ggplot(dat,  
  aes(x = gls)) +  
  geom_histogram(bins = 50) +  
  labs(x = "Goals scored", y = "",  
    subtitle = "Count of NWSL goals per player and year, 2013 - 2019")
```



# Goal scoring

```
ggplot(dat,  
  aes(x = gls)) +  
  geom_histogram(bins = 50) +  
  facet_wrap(~pos, scales = "free_y") +  
  labs(x = "Goals scored", y = "",  
    subtitle = "Count of NWSL goals per player and year, 2013 - 2019")
```



# Modeling goals

```
goals_0<-stan_glm(gls ~ pos,  
  data = dat,  
  family = "poisson",  
  refresh = 0)
```

```
goals_0
```

```
## stan_glm  
## family:      poisson [log]  
## formula:     gls ~ pos  
## observations: 1350  
## predictors:   6  
## -----  
##              Median MAD_SD  
## (Intercept) -1.1    0.1  
## posDF,FW     1.0    0.2  
## posDF,MF     0.9    0.1  
## posFW        2.0    0.1  
## posFW,MF     2.2    0.1  
## posMF        1.2    0.1  
##  
## -----  
## * For help interpreting the printed output see ?print.stanreg  
## * For info on the priors used see ?prior_summary.stanreg
```

## So how many goals does our model expect for each position?

We could just do the math:  $\lambda_i = E(y_i|X) = e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}$

```
exp(coef(goals_0))
```

```
## (Intercept)    posDF,FW    posDF,MF    posFW    posFW,MF    posMF
##      0.335945      2.729285      2.483760      7.284995      9.422848      3.288571
```

And because  $e^{a+b} = e^a \times e^b$

Expected goals for a forward under model 0 are  $e^{\beta_0} \times e^{\beta_3}$

```
# intercept is in row 1, b3 is in row 4
```

```
exp(coef(goals_0)[4]) * exp(coef(goals_0)[1])
```

```
##      posFW
```

```
## 2.447358
```

## So how many goals does our model expect for each position?

Or we could have R handle everything using `predict()`

```
sim_dat <- data.frame(pos = unique(dat$pos))
sim_dat <- sim_dat %>%
  mutate(e_gls = predict(goals_0, newdata = sim_dat, type = "response"))
```

```
sim_dat
```

```
##      pos      e_gls
## 1    DF 0.3367806
## 2 FW,MF 3.1723728
## 3    FW 2.4509501
## 4 DF,MF 0.8365754
## 5    MF 1.1053956
## 6 DF,FW 0.9243009
```

## Regression generates conditional means

Not coincidentally:

```
dat %>% group_by(pos) %>% summarize(gls = mean(gls))
```

```
## # A tibble: 6 x 2
```

```
##   pos      gls
```

```
##   <chr> <dbl>
```

```
## 1 DF      0.334
```

```
## 2 DF,FW  0.925
```

```
## 3 DF,MF  0.837
```

```
## 4 FW      2.45
```

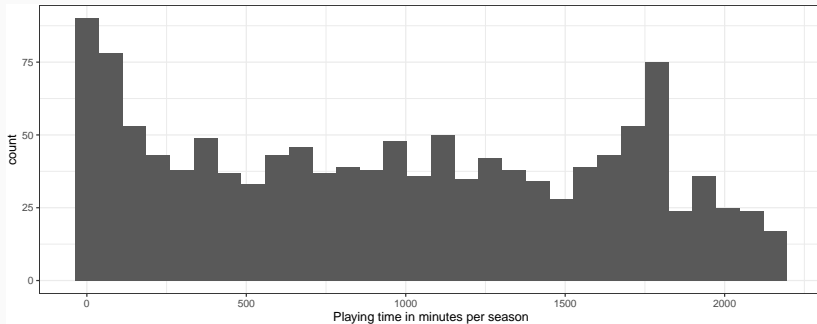
```
## 5 FW,MF  3.17
```

```
## 6 MF      1.11
```

# Fitting a more complex model

Let's look at playing time as a predictor

```
ggplot(dat,  
  aes(x = min)) +  
  geom_histogram() +  
  labs(x = "Playing time in minutes per season")
```



How does playing time impact scoring?



How does playing time impact scoring?

As an opportunity structure - more time = more chances

How does playing time impact scoring?

As an opportunity structure - more time = more chances

Is playing time likely to have the same effect on goal scoring for each position?

How does playing time impact scoring?

As an opportunity structure - more time = more chances

Is playing time likely to have the same effect on goal scoring for each position?

Let's evaluate this model:

$$m1 : E(\text{goals} | \text{position}, \text{minutes}) = e^{\beta_0 + \beta_1 \text{position} + \beta_2 \text{minutes}}$$

# Estimate the model

```
# 1 minute is a small difference, let's use z-scores
goals_1<-stan_glm(gls ~ pos + scale(min),
  data = dat,
  family = "poisson",
  refresh = 0)
```

# Check our fits

```
goals_1
```

```
## stan_glm
## family:      poisson [log]
## formula:      gls ~ pos + scale(min)
## observations: 1271
## predictors:   7
## -----
##              Median MAD_SD
## (Intercept) -1.5    0.1
## posDF,FW     1.3    0.2
## posDF,MF     0.8    0.1
## posFW        2.3    0.1
## posFW,MF     2.4    0.1
## posMF        1.3    0.1
## scale(min)   0.8    0.0
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

What does this show?

## Let's look at model expectations with simulation

```
## all positions with varying play time
pos<-unique(dat$pos)
min<-seq(from = 0,
        to = max(dat$min, na.rm=T),
        by = 5)
sim_dat<-expand_grid(pos, min)

head(sim_dat)
```

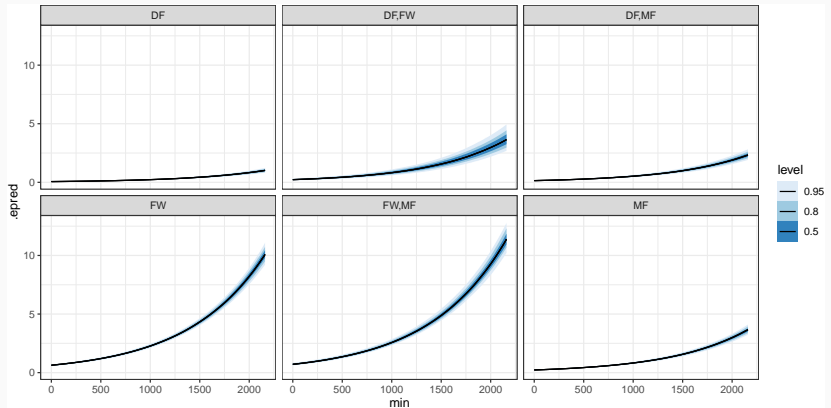
```
## # A tibble: 6 x 2
##   pos      min
##   <chr> <dbl>
## 1 DF      0
## 2 DF      5
## 3 DF     10
## 4 DF     15
## 5 DF     20
## 6 DF     25
```

## Now simulate

```
## add epred draws adds on the expected value scale
library(tidybayes)
sim_dat<-sim_dat %>%
  add_epred_draws(goals_1,
                  ndraws = 500)
```

## Now to visualize expected values

```
ggplot(sim_dat,  
  aes(y = .epred, x = min)) +  
  stat_lineribbon(size = 0.5) + # to make the line thinner  
  facet_wrap(~pos) +  
  scale_fill_brewer()
```






Let's run the simulations again, but now with predicted goals (rather than expected)

```
pos<-unique(dat$pos)
min<-seq(from = 0,
        to = max(dat$min, na.rm=T),
        by = 5)
sim_dat<-expand_grid(pos, min)

sim_dat<-sim_dat %>%
  add_predicted_draws(goals_1,
                    ndraws = 500)
```

# Visualize predicted values

```
ggplot(sim_dat,  
  aes(y = .prediction, x = min)) +  
  stat_lineribbon(size = 0.5) + # to make the line thinner  
  facet_wrap(~pos) +  
  scale_fill_brewer()
```



## Advantages of the Poisson distribution for regression

1. Constrained to non-negative integers
2. Variance scales with the expectation of  $y$  (non-constant error variance!)
3. Relatively simple to interpret